

wjcb86tlq

February 18, 2024

```
[ ]: pip install --upgrade scikit-learn
```

```
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-
packages (1.2.2)
Collecting scikit-learn
  Downloading
scikit_learn-1.4.0-1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
(12.1 MB)

12.1/12.1 MB
35.3 MB/s eta 0:00:00
Requirement already satisfied: numpy<2.0,>=1.19.5 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.23.5)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist-
packages (from scikit-learn) (1.11.4)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-
packages (from scikit-learn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.2.0)
Installing collected packages: scikit-learn
  Attempting uninstall: scikit-learn
    Found existing installation: scikit-learn 1.2.2
    Uninstalling scikit-learn-1.2.2:
      Successfully uninstalled scikit-learn-1.2.2
Successfully installed scikit-learn-1.4.0
```

```
[ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[ ]: data = pd.read_csv("/content/diabetes.csv")
display(data)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	

4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[ ]: data.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[ ]: data.isnull().sum()
```


```
[ ]: Pregnancies          0
      Glucose             0
      BloodPressure       0
      SkinThickness       0
      Insulin             0
      BMI                 0
      DiabetesPedigreeFunction 0
      Age                 0
      Outcome             0
      dtype: int64
```

```
[ ]: data.median()
```

```
[ ]: Pregnancies          3.0000
      Glucose            117.0000
      BloodPressure       72.0000
      SkinThickness       23.0000
      Insulin             30.5000
      BMI                 32.0000
      DiabetesPedigreeFunction 0.3725
      Age                 29.0000
      Outcome             0.0000
      dtype: float64
```


```
[ ]: sns.histplot(data = data, x = 'Pregnancies', kde = True)
```

```
[ ]: <Axes: xlabel='Pregnancies', ylabel='Count'>
```




```
[ ]: sns.histplot(data = data, x = 'Glucose', kde = True)
```

```
[ ]: <Axes: xlabel='Glucose', ylabel='Count'>
```




```
[ ]: sns.histplot(data = data, x = 'BloodPressure', kde = True)
```

```
[ ]: <Axes: xlabel='BloodPressure', ylabel='Count'>
```




```
[ ]: sns.histplot(data = data, x = 'SkinThickness', kde = True)
```

```
[ ]: <Axes: xlabel='SkinThickness', ylabel='Count'>
```




```
[ ]: sns.histplot(data = data, x = 'Insulin', kde = True)
```

```
[ ]: <Axes: xlabel='Insulin', ylabel='Count'>
```




```
[ ]: sns.histplot(data = data, x = 'BMI', kde = True)
```

```
[ ]: <Axes: xlabel='BMI', ylabel='Count'>
```


```
[ ]: sns.histplot(data = data, x = 'DiabetesPedigreeFunction', kde = True)
```

```
[ ]: <Axes: xlabel='DiabetesPedigreeFunction', ylabel='Count'>
```



```
[ ]: sns.histplot(data = data, x = 'Age', kde = True)
```

```
[ ]: <Axes: xlabel='Age', ylabel='Count'>
```



```
[ ]: data[data['Outcome'] == 0].count()
```

```
[ ]: Pregnancies      500
      Glucose          500
      BloodPressure    500
      SkinThickness     500
      Insulin           500
      BMI               500
      DiabetesPedigreeFunction  500
      Age               500
      Outcome           500
      dtype: int64
```

```
[ ]: data[data['Outcome'] == 1].count()
```

```
[ ]: Pregnancies      268
      Glucose          268
      BloodPressure    268
      SkinThickness     268
      Insulin           268
      BMI               268
```

```
DiabetesPedigreeFunction    268
Age                        268
Outcome                    268
dtype: int64
```

```
[ ]: from sklearn.utils import resample

majority_class = data[data['Outcome'] == 0]
minority_class = data[data['Outcome'] == 1]
minority_upsampled = resample(minority_class, replace=True,
    ↳ n_samples=len(majority_class), random_state=42)
data_balanced = pd.concat([majority_class, minority_upsampled], ignore_index =
    ↳ True)
display(data_balanced)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	1	85	66	29	0	26.6
1	1	89	66	23	94	28.1
2	5	116	74	0	0	25.6
3	10	115	0	0	0	35.3
4	4	110	92	0	0	37.6
..
995	7	168	88	42	321	38.2
996	8	143	66	0	0	34.9
997	3	130	78	23	79	28.4
998	6	115	60	39	0	33.7
999	4	184	78	39	277	37.0

	DiabetesPedigreeFunction	Age	Outcome
0	0.351	31	0
1	0.167	21	0
2	0.201	30	0
3	0.134	29	0
4	0.191	30	0
..
995	0.787	40	1
996	0.129	41	1
997	0.323	34	1
998	0.245	40	1
999	0.264	31	1

[1000 rows x 9 columns]

```
[ ]: data_balanced[data_balanced['Outcome'] == 0].count()
```

```
[ ]: Pregnancies    500
      Glucose        500
```

```

BloodPressure      500
SkinThickness      500
Insulin            500
BMI                500
DiabetesPedigreeFunction  500
Age                500
Outcome            500
dtype: int64

```

```
[ ]: data_balanced[data_balanced['Outcome'] == 1].count()
```

```

[ ]: Pregnancies      500
      Glucose          500
      BloodPressure    500
      SkinThickness    500
      Insulin          500
      BMI              500
      DiabetesPedigreeFunction  500
      Age              500
      Outcome          500
dtype: int64

```

```

[ ]: from sklearn.linear_model import LogisticRegression
      from sklearn.preprocessing import StandardScaler, PolynomialFeatures
      from sklearn.model_selection import train_test_split, KFold, cross_val_score,
      ↪ RandomizedSearchCV, ValidationCurveDisplay
      from sklearn.metrics import accuracy_score, precision_score, recall_score,
      ↪ f1_score, classification_report, confusion_matrix

```

```

[ ]: y = data['Outcome']
      y_b = data_balanced['Outcome']
      data = data.drop('Outcome', axis = 1)
      data_balanced = data_balanced.drop('Outcome', axis = 1)

```

```

[ ]: std = StandardScaler()
      data = std.fit_transform(data)
      data_balanced = std.fit_transform(data_balanced)

```

```
[ ]: display(data)
```

```

array([[ 0.63994726,  0.84832379,  0.14964075, ...,  0.20401277,
         0.46849198,  1.4259954 ],
       [-0.84488505, -1.12339636, -0.16054575, ..., -0.68442195,
        -0.36506078, -0.19067191],
       [ 1.23388019,  1.94372388, -0.26394125, ..., -1.10325546,
         0.60439732, -0.10558415],
       ...,

```

```
[ 0.3429808 ,  0.00330087,  0.14964075, ..., -0.73518964,
 -0.68519336, -0.27575966],
 [-0.84488505,  0.1597866 , -0.47073225, ..., -0.24020459,
 -0.37110101,  1.17073215],
 [-0.84488505, -0.8730192 ,  0.04624525, ..., -0.20212881,
 -0.47378505, -0.87137393]])
```

```
[ ]: display(data_balanced)
```

```
array([[ -0.89511603, -1.25513659, -0.21517238, ..., -0.77277546,
        -0.38393058, -0.28171954],
 [ -0.89511603, -1.12702548, -0.21517238, ..., -0.57484782,
        -0.94627816, -1.14245752],
 [  0.22693477, -0.26227547,  0.20102178, ..., -0.90472721,
        -0.84236611, -0.36779334],
 ...,
 [ -0.33409063,  0.18611342,  0.40911886, ..., -0.5352623 ,
        -0.46950522, -0.02349815],
 [  0.50744747, -0.29430325, -0.52731801, ...,  0.16408201,
        -0.70789169,  0.49294464],
 [-0.05357793,  1.91561342,  0.40911886, ...,  0.5995228 ,
        -0.64982319, -0.28171954]])
```

```
[ ]: n = int(input())
poly = PolynomialFeatures(degree = n)
data_poly = poly.fit_transform(data)
data_balanced_poly = poly.fit_transform(data_balanced)
```

2

```
[ ]: print(data_poly.shape)
print(data_balanced_poly.shape)
```

```
(768, 45)
(1000, 45)
```

```
[ ]: display(data_poly)
```

```
array([[ 1.          ,  0.63994726,  0.84832379, ...,  0.21948473,
        0.66806741,  2.03346289],
 [ 1.          , -0.84488505, -1.12339636, ...,  0.13326937,
        0.06960683,  0.03635578],
 [ 1.          ,  1.23388019,  1.94372388, ...,  0.36529612,
        -0.06381478,  0.01114801],
 ...,
 [ 1.          ,  0.3429808 ,  0.00330087, ...,  0.46948994,
        0.18894869,  0.07604339],
 [ 1.          , -0.84488505,  0.1597866 , ...,  0.13771596,
```

```

-0.43445989, 1.37061376],
[ 1.          , -0.84488505, -0.8730192 , ...,  0.22447227,
 0.41284394,  0.75929253]])

```

```
[ ]: display(data_balanced_poly)
```

```

array([[ 1.00000000e+00, -8.95116026e-01, -1.25513659e+00, ...,
         1.47402694e-01,  1.08160748e-01,  7.93658993e-02],
 [ 1.00000000e+00, -8.95116026e-01, -1.12702548e+00, ...,
         8.95442364e-01,  1.08108260e+00,  1.30520918e+00],
 [ 1.00000000e+00,  2.26934774e-01, -2.62275473e-01, ...,
         7.09580666e-01,  3.09816644e-01,  1.35271939e-01],
 ...,
 [ 1.00000000e+00, -3.34090626e-01,  1.86113417e-01, ...,
         2.20435148e-01,  1.10325025e-02,  5.52162903e-04],
 [ 1.00000000e+00,  5.07447475e-01, -2.94303251e-01, ...,
         5.01110645e-01, -3.48951414e-01,  2.42994418e-01],
 [ 1.00000000e+00, -5.35779257e-02,  1.91561342e+00, ...,
         4.22270179e-01,  1.83067890e-01,  7.93658993e-02]])

```

```
[ ]: data = np.tile(data, (10,1))
data_balanced = np.tile(data_balanced, (10,1))
data_poly = np.tile(data_poly, (10,1))
data_balanced_poly = np.tile(data_balanced_poly, (10,1))
y = np.tile(y, 10)
y_b = np.tile(y_b, 10)
```

```
[ ]: x_train, x_temp, y_train, y_temp = train_test_split(data, y, test_size=0.4,
    ↪random_state=42, stratify = y)
x_val, x_test, y_val, y_test = train_test_split(x_temp, y_temp, test_size=0.5,
    ↪random_state=42, stratify = y_temp)
x_train_b, x_temp_b, y_train_b, y_temp_b = train_test_split(data_balanced, y_b,
    ↪test_size=0.4, random_state=42, stratify = y_b)
x_val_b, x_test_b, y_val_b, y_test_b = train_test_split(x_temp_b, y_temp_b,
    ↪test_size=0.5, random_state=42, stratify = y_temp_b) #b for balanced
x_train_p, x_temp_p, y_train_p, y_temp_p = train_test_split(data_poly, y,
    ↪test_size=0.4, random_state=42, stratify = y)
x_val_p, x_test_p, y_val_p, y_test_p = train_test_split(x_temp_p, y_temp_p,
    ↪test_size=0.5, random_state=42, stratify = y_temp_p)
x_train_b_p, x_temp_b_p, y_train_b_p, y_temp_b_p =
    ↪train_test_split(data_balanced_poly, y_b, test_size=0.4, random_state=42,
    ↪stratify = y_b)
x_val_b_p, x_test_b_p, y_val_b_p, y_test_b_p = train_test_split(x_temp_b_p,
    ↪y_temp_b_p, test_size=0.5, random_state=42, stratify = y_temp_b_p) #p for
    ↪polygonal features
```

```
[ ]: print(x_train.shape, x_test.shape, x_val.shape)
      print(x_train_b.shape, x_test_b.shape, x_val_b.shape)
      print(x_train_p.shape, x_test_p.shape, x_val_p.shape)
      print(x_train_b_p.shape, x_test_b_p.shape, x_val_b_p.shape)
```

```
(460, 8) (154, 8) (154, 8)
(600, 8) (200, 8) (200, 8)
(460, 45) (154, 45) (154, 45)
(600, 45) (200, 45) (200, 45)
```

```
[ ]: model_logreg = LogisticRegression(random_state = 42, max_iter = 10000)
```

```
[ ]: def kfold(input_data, output_data):
      kf = KFold( n_splits = int(input_data.shape[0]/10), shuffle = True,
      random_state = 42)
      scores = cross_val_score(model_logreg, input_data, output_data, cv = kf)
      print("Cross validation scores : ", scores)
      print("Mean cross validation score : ", np.mean(scores))
```

```
[ ]: kfold(x_train, y_train)
      kfold(x_train_b, y_train_b)
      kfold(x_train_p, y_train_p)
      kfold(x_train_b_p, y_train_b_p)
```

```
Cross validation scores : [1.  0.7 0.7 0.8 0.9 0.7 0.6 0.7 0.9 0.8 0.8 0.8 0.9
0.9 0.7 0.9 0.9 0.7
 0.9 0.8 0.8 0.8 0.7 0.5 0.8 0.5 0.6 0.7 0.8 1.  0.9 0.6 0.7 0.8 1.  1.
 0.8 1.  0.8 0.9 0.7 0.8 0.8 0.6 0.8 0.7]
Mean cross validation score : 0.7869565217391304
Cross validation scores : [0.8 0.4 0.8 0.7 0.6 0.6 0.7 0.8 0.8 0.7 0.6 0.7 0.9
1.  0.7 0.7 0.8 0.8
 0.8 0.6 0.7 0.8 0.7 0.9 0.6 0.6 0.7 0.8 0.5 0.6 0.6 0.6 0.7 0.8 0.6 0.8
 0.9 0.8 0.7 0.9 0.7 0.8 0.6 0.9 0.7 0.8 0.7 0.9 0.9 0.7 0.7 0.9 0.8 0.4
 0.6 0.9 0.6 0.9 0.8 0.9]
Mean cross validation score : 0.7333333333333333
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-425-391bbad65a22> in <cell line: 3>()
      1 kfold(x_train, y_train)
      2 kfold(x_train_b, y_train_b)
----> 3 kfold(x_train_p, y_train_p)
      4 kfold(x_train_b_p, y_train_b_p)

<ipython-input-424-50f5deeac830> in kfold(input_data, output_data)
      1 def kfold(input_data, output_data):
```



```

2     kf = KFold( n_splits = int(input_data.shape[0]/10), shuffle = True,
↳random_state = 42)
----> 3     scores = cross_val_score(model_logreg, input_data, output_data, cv =
↳kf)
4     print("Cross validation scores : ", scores)
5     print("Mean cross validation score : ", np.mean(scores))

/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py in
↳wrapper(*args, **kwargs)
211         )
212     ):
--> 213         return func(*args, **kwargs)
214     except InvalidParameterError as e:
215         # When the function is just a wrapper around an
↳estimator, we allow

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py
↳in cross_val_score(estimator, X, y, groups, scoring, cv, n_jobs, verbose,
↳fit_params, params, pre_dispatch, error_score)
712     scorer = check_scoring(estimator, scoring=scoring)
713
--> 714     cv_results = cross_validate(
715         estimator=estimator,
716         X=X,

/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py in
↳wrapper(*args, **kwargs)
211         )
212     ):
--> 213         return func(*args, **kwargs)
214     except InvalidParameterError as e:
215         # When the function is just a wrapper around an
↳estimator, we allow

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py
↳in cross_validate(estimator, X, y, groups, scoring, cv, n_jobs, verbose,
↳fit_params, params, pre_dispatch, return_train_score, return_estimator,
↳return_indices, error_score)
423     # independent, and that it is pickle-able.
424     parallel = Parallel(n_jobs=n_jobs, verbose=verbose,
↳pre_dispatch=pre_dispatch)
--> 425     results = parallel(
426         delayed(_fit_and_score)(
427             clone(estimator),

/usr/local/lib/python3.10/dist-packages/sklearn/utils/parallel.py in
↳__call__(self, iterable)

```

```

        65             for delayed_func, args, kwargs in iterable
        66         )
--> 67         return super().__call__(iterable_with_config)
        68
        69

/usr/local/lib/python3.10/dist-packages/joblib/parallel.py in __call__(self,
↳iterable)
    1861         output = self._get_sequential_output(iterable)
    1862         next(output)
-> 1863         return output if self.return_generator else list(output)
    1864
    1865         # Let's create an ID that uniquely identifies the current call.
↳If the

/usr/local/lib/python3.10/dist-packages/joblib/parallel.py in
↳_get_sequential_output(self, iterable)
    1790         self.n_dispatched_batches += 1
    1791         self.n_dispatched_tasks += 1
-> 1792         res = func(*args, **kwargs)
    1793         self.n_completed_tasks += 1
    1794         self.print_progress()

/usr/local/lib/python3.10/dist-packages/sklearn/utils/parallel.py in
↳__call__(self, *args, **kwargs)
    127         config = {}
    128         with config_context(**config):
--> 129         return self.function(*args, **kwargs)

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py
↳in_fit_and_score(estimator, X, y, scorer, train, test, verbose, parameters,
↳fit_params, score_params, return_train_score, return_parameters,
↳return_n_test_samples, return_times, return_estimator, split_progress,
↳candidate_progress, error_score)
    888         estimator.fit(X_train, **fit_params)
    889     else:
--> 890         estimator.fit(X_train, y_train, **fit_params)
    891
    892     except Exception:

/usr/local/lib/python3.10/dist-packages/sklearn/base.py in wrapper(estimator,
↳*args, **kwargs)
    1349         )
    1350     ):
-> 1351         return fit_method(estimator, *args, **kwargs)
    1352
    1353     return wrapper

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py in
↳fit(self, X, y, sample_weight)
    1294         n_threads = 1
    1295
-> 1296         fold_coefs_ = Parallel(n_jobs=self.n_jobs, verbose=self.verbose,
↳prefer=prefer)(
    1297             path_func(
    1298                 X,

/usr/local/lib/python3.10/dist-packages/sklearn/utils/parallel.py in
↳__call__(self, iterable)
    65         for delayed_func, args, kwargs in iterable
    66     )
---> 67     return super().__call__(iterable_with_config)
    68
    69

/usr/local/lib/python3.10/dist-packages/joblib/parallel.py in __call__(self,
↳iterable)
    1861         output = self._get_sequential_output(iterable)
    1862         next(output)
-> 1863         return output if self.return_generator else list(output)
    1864
    1865         # Let's create an ID that uniquely identifies the current call.
↳If the

/usr/local/lib/python3.10/dist-packages/joblib/parallel.py in
↳_get_sequential_output(self, iterable)
    1790         self.n_dispatched_batches += 1
    1791         self.n_dispatched_tasks += 1
-> 1792         res = func(*args, **kwargs)
    1793         self.n_completed_tasks += 1
    1794         self.print_progress()

/usr/local/lib/python3.10/dist-packages/sklearn/utils/parallel.py in
↳__call__(self, *args, **kwargs)
    127         config = {}
    128         with config_context(**config):
--> 129         return self.function(*args, **kwargs)

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py in
↳_logistic_regression_path(X, y, pos_class, Cs, fit_intercept, max_iter, tol,
↳verbose, solver, coef, class_weight, dual, penalty, intercept_scaling,
↳multi_class, random_state, check_input, max_squared_sum, sample_weight,
↳l1_ratio, n_threads)
    453         np.searchsorted(np.array([0, 1, 2, 3]), verbose)
    454     ]

```

```

--> 455             opt_res = optimize.minimize(
456                 func,
457                 w0,

/usr/local/lib/python3.10/dist-packages/scipy/optimize/_minimize.py in
↳ minimize(fun, x0, args, method, jac, hess, hessp, bounds, constraints, tol,
↳ callback, options)
708                 **options)
709     elif meth == 'l-bfgs-b':
--> 710         res = _minimize_lbfgsb(fun, x0, args, jac, bounds,

711                                 callback=callback, **options)
712     elif meth == 'tnc':

/usr/local/lib/python3.10/dist-packages/scipy/optimize/_lbfgsb_py.py in
↳ _minimize_lbfgsb(fun, x0, args, jac, bounds, disp, maxcor, ftol, gtol, eps,
↳ maxfun, maxiter, iprint, callback, maxls, finite_diff_rel_step,
↳ **unknown_options)
363         # until the completion of the current minimization iteratio.
364         # Overwrite f and g:
--> 365         f, g = func_and_grad(x)
366         elif task_str.startswith(b'NEW_X'):
367             # new iteration

/usr/local/lib/python3.10/dist-packages/scipy/optimize/_differentiable_function.
↳ py in fun_and_grad(self, x)
283         if not np.array_equal(x, self.x):
284             self._update_x_impl(x)
--> 285         self._update_fun()
286         self._update_grad()
287         return self.f, self.g

/usr/local/lib/python3.10/dist-packages/scipy/optimize/_differentiable_function.
↳ py in _update_fun(self)
249     def _update_fun(self):
250         if not self.f_updated:
--> 251             self._update_fun_impl()
252             self.f_updated = True
253

/usr/local/lib/python3.10/dist-packages/scipy/optimize/_differentiable_function.
↳ py in update_fun()
153
154     def update_fun():
--> 155         self.f = fun_wrapped(self.x)
156
157         self._update_fun_impl = update_fun

```

```

/usr/local/lib/python3.10/dist-packages/scipy/optimize/_differentiable_function.
↳py in fun_wrapped(x)
    135         # Overwriting results in undefined behaviour because
    136         # fun(self.x) will change self.x, with the two no longer
↳linked.
--> 137         fx = fun(np.copy(x), *args)
    138         # Make sure the function returns a true scalar
    139         if not np.isscalar(fx):

/usr/local/lib/python3.10/dist-packages/scipy/optimize/_optimize.py in
↳__call__(self, x, *args)
    75     def __call__(self, x, *args):
    76         """ returns the function value """
---> 77         self._compute_if_needed(x, *args)
    78         return self._value
    79

/usr/local/lib/python3.10/dist-packages/scipy/optimize/_optimize.py in
↳_compute_if_needed(self, x, *args)
    69         if not np.all(x == self.x) or self._value is None or self.jac is
↳None:
    70             self.x = np.asarray(x).copy()
---> 71             fg = self.fun(x, *args)
    72             self.jac = fg[1]
    73             self._value = fg[0]

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_linear_loss.py in
↳loss_gradient(self, coef, X, y, sample_weight, l2_reg_strength, n_threads,
↳raw_prediction)
    274
    275     if raw_prediction is None:
--> 276         weights, intercept, raw_prediction = self.
↳weight_intercept_raw(coef, X)
    277     else:
    278         weights, intercept = self.weight_intercept(coef)

KeyboardInterrupt:

```

```

[ ]: model_logreg_ = model_logreg.fit(x_train, y_train)
y_pred_ = model_logreg_.predict(x_val)
print(accuracy_score(y_val, y_pred_))
y_pred_ = model_logreg_.predict(x_test)
print(accuracy_score(y_test, y_pred_))

```

```

0.7337662337662337
0.7402597402597403

```

```
[ ]: model_logreg_b = model_logreg.fit(x_train_b, y_train_b)
      y_pred_b = model_logreg_b.predict(x_val_b)
      print(accuracy_score(y_val_b, y_pred_b))
      y_pred_b = model_logreg_b.predict(x_test_b)
      print(accuracy_score(y_test_b, y_pred_b))
```

0.71

0.715

```
[ ]: model_logreg_p = model_logreg.fit(x_train_p, y_train_p)
      y_pred_p = model_logreg_p.predict(x_val_p)
      print(accuracy_score(y_val_p, y_pred_p))
      y_pred_p = model_logreg_p.predict(x_test_p)
      print(accuracy_score(y_test_p, y_pred_p))
```

0.7857142857142857

0.6753246753246753

```
[ ]: model_logreg_b_p = model_logreg.fit(x_train_b_p, y_train_b_p)
      y_pred_b_p = model_logreg_p.predict(x_val_b_p)
      print(accuracy_score(y_val_b_p, y_pred_b_p))
      y_pred_b_p = model_logreg_p.predict(x_test_b_p)
      print(accuracy_score(y_test_b_p, y_pred_b_p))
```

0.78

0.71

```
[ ]:
```